

# Working with Images

Reading, Exploring and Analyzing, Feature Extraction



**Yordan Darakchiev**  
Technical Trainer



**SoftUni**



Software University

<https://softuni.bg>

# Have a Question?

[sli.do](https://sli.do)

**#DataScience**

- Image processing
  - Reading
  - Exploring
  - Manipulation
  - Convolution
  - Image Morphology





# **Image Processing**

## Understanding What People See

- There are many ways to **read an image**
  - One of the easiest is using scikit-image

```
from skimage.io import imread  
tiger_image = imread("tiger.jpg")
```


- Displaying the image

```
plt.imshow(tiger_image)
```

- The image is actually a **matrix of pixels**
  - Each pixel is an array of three values: **R, G, B  $\in [0; 255]$**
  - Grayscale images only have **one value per pixel**
- Most image **processing algorithms** are easier to understand on grayscale images

```
red = tiger_image[:, :, 0]  
green = tiger_image[:, :, 1]  
blue = tiger_image[:, :, 2]
```

- As usual, histograms tell us **how the values are distributed**
  - How many dark values, how many light values
  - Maximum brightness, peaks, etc.



- Histograms need to have a **single variable**
  - Take each channel separately, e.g., red
  - Convert the **2D matrix to 1D array**: `image.ravel()`
  - Show the histogram as usual
    - It's common to use 256 bins

```
plt.hist(red.ravel(), bins = 256, color = "red")  
plt.show()
```

- We can also plot **all channels on a single histogram**



- Sometimes working **per channel is not necessary**
  - We can combine **all three channels** and get a grayscale image
  - Simplest way: **get the mean of all values**

```
tiger_grayscale = np.mean(tiger_image, axis = 2)
```

- Better way: use coefficients for each channel

- The human eye **discerns colors differently**
- Were more sensitive to green colors
- Some formulas are given [here](#)

```
tiger_grayscale = 0.299 * red + 0.587 * green + 0.114 * blue
```

- Depending on the image, the differences **may or may not be easy to see**
- For art purposes, we can experiment with our own **coefficients for combining all channels**

- **Convolution kernel (filter)**
  - A small, usually 3x3, matrix of numbers
- Convolution process
  - Input: image, kernel
  - Output: new image

35	40	41	45	50
40	40	42	46	52
42	46	50	55	55
48	52	56	58	60
56	60	65	70	75

×

	0	1	0	
	0	0	0	
	0	0	0	

=

		42		

Combining the image and a kernel:

- Apply the kernel **over each pixel**
- Multiply the values **element-wise** (Hadamard product)
- Sum **all values**
- Assign the **sum to the corresponding pixel** in the output image
  - Image corners are treated in different ways, not really important how

35	40	41	45	50
40	40	42	46	52
42	46	50	55	55
48	52	56	58	60
56	60	65	70	75

×


	0	1	0	
	0	0	0	
	0	0	0	

=

		42		

- The choice of kernel depends what the **output image will represent**


```
from scipy.ndimage.filters import convolve  
convolve(image, kernel)
```



## Example: box blur


```
box_blur_kernel = np.array([  
    [1, 1, 1],  
    [1, 1, 1],  
    [1, 1, 1]  
]) / 9
```

```
blurred = convolve(tiger_grayscale, box_blur_kernel)  
plt.imshow(tiger_grayscale[150:250, 300:400], cmap = "gray")  
plt.show()  
plt.imshow(blurred[150:250, 300:400], cmap = "gray")  
plt.show()
```



- Four main operations (see [this](#) tutorial)
  - **Dilation**
  - **Erosion**
  - **Opening**
  - **Closing**
- A simple series of **algorithms for image transformation**
- Basic methodology
  - Choose a structuring element (e.g., 2x2 square or cross)
  - Move the element around the image
  - Apply an operation

- Input: **binary image**
  - Pixel values **0** and **1**, not **[0; 255]**
  - This is called **thresholding**
- Output: **transformed image**






First get **all values** inside the structuring element


- **Erosion**: replace all values with the min value

- Strips away a layer of pixels
- Holes become larger
- Small regions are eliminated





- **Dilation**: replace all values with the max value

- Adds a layer of pixels
- Gaps become smaller
- Small gaps are filled in



- **Opening**: erosion followed by dilation
  - Pixels which survived erosion are restored to their original size
  - Opens up a gap between two objects connected by thin bridges
- **Closing**: dilation followed by erosion
  - Fills in holes in the regions while keeping the initial region sizes



- Matrix operations – **pixel-wise**
  - One image:
    - **Addition, Gain, Negative**
    - **Resampling, Cutting**
  - Geometric transformations – **perspective, warp**, etc.
  - Two (or more) images:
    - **Addition (multiple exposure)**
    - **Subtraction (difference)**
    - **Division (normalization)**
    - **Averaging**


- Thresholding (usually 2 levels)
- Fourier transform, filtering and convolution
- Contrast enhancement
- Histogram equalization
- Stacking (many 2D images  $\Rightarrow$  one 3D image)
- Analysis:
  - Measurements, Segmentation, Object extraction / Identification
  - Enhancements, Inpainting

# Summary

- Image processing
  - Reading
  - Exploring
  - Manipulation
  - Convolution
  - Image Morphology



# Questions?



# SoftUni Diamond Partners



- Software University – High-Quality Education, Profession and Job for Software Developers
  - [softuni.bg](http://softuni.bg), [about.softuni.bg](http://about.softuni.bg)
- Software University Foundation
  - [softuni.foundation](http://softuni.foundation)
- Software University @ Facebook
  - [facebook.com/SoftwareUniversity](https://facebook.com/SoftwareUniversity)





- This course (slides, examples, demos, exercises, homework, documents, videos and other assets) is **copyrighted content**
- Unauthorized copy, reproduction or use is illegal
- © SoftUni – <https://about.softuni.bg/>
- © Software University – <https://softuni.bg>

